BIOS6301: Homework 5

Sarah Lotspeich

15 November 2016

Question 1

24 points

Import the HAART dataset (haart.csv) from the GitHub repository into R, and perform the following manipulations: (4 points each)

```r
library(lubridate)
haart <- read.csv("https://raw.githubusercontent.com/fonnesbeck/Bios6301/master/datasets/haart.csv")
head(haart) #preview structure of base dataset
```

```
##   male age aids cd4baseline logvl  weight hemoglobin    init.reg init.date
## 1    1  25    0          NA    NA      NA         NA 3TC,AZT,EFV    7/1/03
## 2    1  49    0         143    NA 58.0608         11 3TC,AZT,EFV  11/23/04
## 3    1  42    1         102    NA 48.0816          1 3TC,AZT,EFV   4/30/03
## 4    0  33    0         107    NA 46.0000         NA 3TC,AZT,NVP   3/25/06
## 5    1  27    0          52     4      NA         NA 3TC,D4T,EFV    9/1/04
## 6    0  34    0         157    NA 54.8856         NA 3TC,AZT,NVP   12/2/03
##   last.visit death date.death
## 1    2/26/07     0        <NA>
## 2    2/22/08     0        <NA>
## 3   11/21/05     1     1/11/06
## 4     5/5/06     1      5/7/06
## 5   11/13/07     0        <NA>
## 6    2/28/08     0        <NA>
```

1. Convert date columns into a usable (for analysis) format. Use the table command to display the counts of the year from init.date.

```r
# convert date columns
haart[, "init.date"] <- as.POSIXct(haart[, "init.date"], format = "%m/%d/%y")
haart[, "last.visit"] <- as.POSIXct(haart[, "last.visit"], format = "%m/%d/%y")
haart[, "date.death"] <- as.POSIXct(haart[, "date.death"], format = "%m/%d/%y")

# display counts of the year from init.date
table(year(haart[, "init.date"]))
```

```
##
## 1998 2000 2001 2002 2003 2004 2005 2006 2007
##    1    5   17   60  270  292  207  104   44
```

1.2. Create an indicator variable (one which takes the values 0 or 1 only) to represent death within 1 year of the initial visit. How many observations died in year 1?

```
# create indicator variable to represent death within 1 year of the initial
# visit
haart$death.within.year <- (difftime(haart$date.death, haart$last.visit, units = "days") <=
    365)
haart$death.within.year[is.na(haart$death.within.year)] <- 0
head(haart)  #prove I added the column correctly
```

```
##   male age aids cd4baseline logvl  weight hemoglobin    init.reg
## 1    1  25    0          NA    NA      NA         NA 3TC,AZT,EFV
## 2    1  49    0         143    NA 58.0608         11 3TC,AZT,EFV
## 3    1  42    1         102    NA 48.0816          1 3TC,AZT,EFV
## 4    0  33    0         107    NA 46.0000         NA 3TC,AZT,NVP
## 5    1  27    0          52     4      NA         NA 3TC,D4T,EFV
## 6    0  34    0         157    NA 54.8856         NA 3TC,AZT,NVP
##    init.date last.visit death date.death death.within.year
## 1 2003-07-01 2007-02-26     0       <NA>                 0
## 2 2004-11-23 2008-02-22     0       <NA>                 0
## 3 2003-04-30 2005-11-21     1 2006-01-11                 1
## 4 2006-03-25 2006-05-05     1 2006-05-07                 1
## 5 2004-09-01 2007-11-13     0       <NA>                 0
## 6 2003-12-02 2008-02-28     0       <NA>                 0
```

```
# calculate number of paricipants who died in year 1
(yr1.deaths <- sum(haart$death.within.year))
```

```
## [1] 104
```

1.3. Use the init.date, last.visit and death.date columns to calculate a followup time (in days), which is the difference between the first and either the last visit or a death event (whichever comes first).

```
#function that takes who dataframe, adds a followup time column and calculates values
#then returns the new and improved dataframe
calculate.followup <- function(dataframe)
{
  for (i in 1:nrow(dataframe))
  {
    if (dataframe$death[i] == 1) #check to see if death came first
    {
      dataframe$followup.time[i] <- difftime(dataframe$date.death[i],dataframe$init.date[i],
                                      units="days")
    }
    else
    {
      dataframe$followup.time[i] <- difftime(dataframe$last.visit[i],dataframe$init.date[i],
                                      units="days")
    }
  }
  return(dataframe)
}

haart <- calculate.followup(haart)
head(haart)
```

```
##   male age aids cd4baseline logvl  weight hemoglobin    init.reg
## 1    1  25    0          NA    NA      NA         NA 3TC,AZT,EFV
## 2    1  49    0         143    NA 58.0608         11 3TC,AZT,EFV
## 3    1  42    1         102    NA 48.0816          1 3TC,AZT,EFV
## 4    0  33    0         107    NA 46.0000         NA 3TC,AZT,NVP
## 5    1  27    0          52     4      NA         NA 3TC,D4T,EFV
## 6    0  34    0         157    NA 54.8856         NA 3TC,AZT,NVP
##     init.date last.visit death date.death death.within.year followup.time
## 1 2003-07-01 2007-02-26     0       <NA>                 0    1336.04167
## 2 2004-11-23 2008-02-22     0       <NA>                 0    1186.00000
## 3 2003-04-30 2005-11-21     1 2006-01-11                 1     987.04167
## 4 2006-03-25 2006-05-05     1 2006-05-07                 1      42.95833
## 5 2004-09-01 2007-11-13     0       <NA>                 0    1168.04167
## 6 2003-12-02 2008-02-28     0       <NA>                 0    1549.00000
```

1.4. If these times are longer than 1 year, censor them (this means if the value is above 365, set followup to 365).

```
#function that takes in a vector of followup.times and
#censors them so that 365 is the maximum, then returns the censored vector
max(haart$followup.time) #see the max of followup.time as calculated
```

```
## [1] 3533.042
```

```
censor.followup <- function(followup.time)
{
  for (i in 1:length(followup.time))
  {
    if (followup.time[i] > 365)
    {
      followup.time[i] <- 365
    }
  }
  return(followup.time)
}
haart$followup.time <- censor.followup(haart$followup.time)
max(haart$followup.time) #see new max of followup.time after censoring
```

```
## [1] 365
```

1.5. Print the quantile for this new variable.

```
quantile(haart$followup.time)
```

```
##    0%   25%   50%   75%  100%
##   0.0 329.5 365.0 365.0 365.0
```

1.6. Create another indicator variable representing loss to followup; this means the observation is not known to be dead but does not have any followup visits after the first year. How many records are lost-to-followup?

```r
#function that takes in whole dataframe and appends an indicator
#column for whether or not individuals were lost fo followup
calculate.losstofollowup <- function(dataframe)
{
  for (i in 1:nrow(dataframe))
  {
    if (dataframe$death[i] == 1)#if dead, then they weren't "lost-to-followup"
    {
      dataframe$loss.to.followup[i] <- 0
    }
    else
    {
      #if not dead, haven't had followup after first year, "lost-to-followup"
      if (as.integer(difftime(dataframe$last.visit[i],dataframe$init.date[i],units="days")) <= 365)
      {
        dataframe$loss.to.followup[i] <- 1
      }
      #if not dead, have had a followup after first year, not "lost-to-followup"
      else dataframe$loss.to.followup[i] <- 0
    }
  }
  return(dataframe)
}


#add lost-to-followup column to the haart dataframe
haart <- calculate.losstofollowup(haart)
head(haart) #appreciate shiny new column
```

```
##   male age aids cd4baseline logvl  weight hemoglobin    init.reg
## 1    1  25    0          NA    NA      NA         NA 3TC,AZT,EFV
## 2    1  49    0         143    NA 58.0608         11 3TC,AZT,EFV
## 3    1  42    1         102    NA 48.0816          1 3TC,AZT,EFV
## 4    0  33    0         107    NA 46.0000         NA 3TC,AZT,NVP
## 5    1  27    0          52     4      NA         NA 3TC,D4T,EFV
## 6    0  34    0         157    NA 54.8856         NA 3TC,AZT,NVP
##    init.date last.visit death date.death death.within.year followup.time
## 1 2003-07-01 2007-02-26     0       <NA>                 0     365.00000
## 2 2004-11-23 2008-02-22     0       <NA>                 0     365.00000
## 3 2003-04-30 2005-11-21     1 2006-01-11                 1     365.00000
## 4 2006-03-25 2006-05-05     1 2006-05-07                 1      42.95833
## 5 2004-09-01 2007-11-13     0       <NA>                 0     365.00000
## 6 2003-12-02 2008-02-28     0       <NA>                 0     365.00000
##   loss.to.followup
## 1                0
## 2                0
## 3                0
## 4                0
## 5                0
## 6                0
```

```r
#sum of indicator variables gives number of patients "lost-to-followup"
sum(haart$loss.to.followup)
```

```
## [1] 173
```

So, from this we can see that 173 records were lost-to-followup.

1.7. Recall our work in class, which separated the init.reg field into a set of indicator variables, one for each unique drug. Create these fields and append them to the database as new columns.

```
create.regimens <- function(dataframe) {
    init.reg <- as.character(dataframe[, "init.reg"])
    dataframe[["init.reg_list"]] <- strsplit(init.reg, ",")
    unique.drugs <- unique(unlist(dataframe$init.reg_list))
    reg.drugs <- matrix(FALSE, nrow = nrow(dataframe), ncol = length(unique.drugs))
    for (i in seq_along(unique.drugs)) {
        reg.drugs[, i] <- sapply(dataframe$init.reg_list, function(x) unique.drugs[i] %in%
            x)
    }
    reg.drugs <- data.frame(reg.drugs)
    names(reg.drugs) <- unique.drugs
    dataframe <- cbind(dataframe, reg.drugs)
    return(dataframe)
}

haart <- create.regimens(haart)
head(haart)   #why not
```

```
##   male age aids cd4baseline logvl  weight hemoglobin    init.reg
## 1    1  25    0          NA    NA      NA         NA 3TC,AZT,EFV
## 2    1  49    0         143    NA 58.0608         11 3TC,AZT,EFV
## 3    1  42    1         102    NA 48.0816          1 3TC,AZT,EFV
## 4    0  33    0         107    NA 46.0000         NA 3TC,AZT,NVP
## 5    1  27    0          52     4      NA         NA 3TC,D4T,EFV
## 6    0  34    0         157    NA 54.8856         NA 3TC,AZT,NVP
##    init.date last.visit death date.death death.within.year followup.time
## 1 2003-07-01 2007-02-26     0       <NA>                 0     365.00000
## 2 2004-11-23 2008-02-22     0       <NA>                 0     365.00000
## 3 2003-04-30 2005-11-21     1 2006-01-11                 1     365.00000
## 4 2006-03-25 2006-05-05     1 2006-05-07                 1      42.95833
## 5 2004-09-01 2007-11-13     0       <NA>                 0     365.00000
## 6 2003-12-02 2008-02-28     0       <NA>                 0     365.00000
##   loss.to.followup init.reg_list  3TC   AZT   EFV   NVP   D4T   ABC   DDI
## 1                0 3TC, AZT, EFV TRUE  TRUE  TRUE FALSE FALSE FALSE FALSE
## 2                0 3TC, AZT, EFV TRUE  TRUE  TRUE FALSE FALSE FALSE FALSE
## 3                0 3TC, AZT, EFV TRUE  TRUE  TRUE FALSE FALSE FALSE FALSE
## 4                0 3TC, AZT, NVP TRUE  TRUE FALSE  TRUE FALSE FALSE FALSE
## 5                0 3TC, D4T, EFV TRUE FALSE  TRUE FALSE  TRUE FALSE FALSE
## 6                0 3TC, AZT, NVP TRUE  TRUE FALSE  TRUE FALSE FALSE FALSE
##     IDV   LPV   RTV   SQV   FTC   TDF   DDC   NFV   T20   ATV   FPV
## 1 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 2 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 3 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 4 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 5 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 6 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

1.8. Which drug regimen are found over 100 times?

```r
regimen <- matrix(nrow = nrow(haart), ncol = 1)
for (j in 17:34) {
    for (i in 1:nrow(haart)) {
        if (haart[i, j] == TRUE) {
            if (j == 17) {
                regimen[i] <- colnames(haart)[j]
            } else {
                regimen[i] <- paste(regimen[i], colnames(haart)[j])
            }
        }
    }
}
haart <- cbind(haart, regimen)
haart[, 17:34] <- NULL  #get rid of the indicators before preview
head(haart)  #check what that actually did
```

```
##   male age aids cd4baseline logvl  weight hemoglobin   init.reg
## 1    1  25    0          NA    NA      NA         NA 3TC,AZT,EFV
## 2    1  49    0         143    NA 58.0608         11 3TC,AZT,EFV
## 3    1  42    1         102    NA 48.0816          1 3TC,AZT,EFV
## 4    0  33    0         107    NA 46.0000         NA 3TC,AZT,NVP
## 5    1  27    0          52     4      NA         NA 3TC,D4T,EFV
## 6    0  34    0         157    NA 54.8856         NA 3TC,AZT,NVP
##    init.date last.visit death date.death death.within.year followup.time
## 1 2003-07-01 2007-02-26     0       <NA>                 0     365.00000
## 2 2004-11-23 2008-02-22     0       <NA>                 0     365.00000
## 3 2003-04-30 2005-11-21     1 2006-01-11                 1     365.00000
## 4 2006-03-25 2006-05-05     1 2006-05-07                 1      42.95833
## 5 2004-09-01 2007-11-13     0       <NA>                 0     365.00000
## 6 2003-12-02 2008-02-28     0       <NA>                 0     365.00000
##   loss.to.followup init.reg_list     regimen
## 1                0 3TC, AZT, EFV 3TC AZT EFV
## 2                0 3TC, AZT, EFV 3TC AZT EFV
## 3                0 3TC, AZT, EFV 3TC AZT EFV
## 4                0 3TC, AZT, NVP 3TC AZT NVP
## 5                0 3TC, D4T, EFV 3TC EFV D4T
## 6                0 3TC, AZT, NVP 3TC AZT NVP
```

```r
sort(table(regimen))
```

```
## regimen
##       3TC ABC IDV RTV          3TC ABC RTV   3TC AZT ABC LPV RTV
##                     1                    1                     1
##    3TC AZT ABC RTV SQV          3TC AZT DDI      3TC AZT EFV NFV
##                     1                    1                     1
##       3TC AZT RTV FPV          3TC EFV TDF      3TC LPV RTV TDF
##                     1                    1                     1
##       3TC RTV TDF FPV   NA ABC DDI LPV RTV   NA ABC DDI RTV ATV
##                     1                    1                     1
##    NA D4T ABC LPV RTV   NA D4T ABC RTV SQV      NA D4T RTV SQV
##                     1                    1                     1
```

6

```
## NA DDI LPV RTV SQV T20           NA EFV D4T DDC           NA EFV DDI FTC
##                        1                        1                        1
##          NA NVP FTC TDF     NA RTV FTC TDF ATV         3TC D4T LPV RTV
##                        1                        1                        2
##             3TC NVP ABC           NA AZT EFV DDI           NA EFV D4T ABC
##                        2                        2                        2
##      NA LPV RTV FTC TDF         NA NVP D4T DDI           NA NVP LPV RTV
##                        2                        2                        2
##             3TC D4T NFV           NA EFV FTC TDF         3TC ABC RTV SQV
##                        3                        3                        4
##             3TC AZT NFV         3TC DDI LPV RTV            NA EFV D4T DDI
##                        4                        4                        4
##         3TC D4T IDV RTV            3TC NVP DDI         3TC AZT IDV RTV
##                        6                        6                        8
##         3TC D4T RTV SQV            3TC EFV ABC            3TC AZT IDV
##                        8                       11                       12
##         3TC AZT RTV SQV            3TC EFV DDI         3TC AZT LPV RTV
##                       13                       15                       16
##             3TC AZT ABC            3TC EFV D4T            3TC NVP D4T
##                       29                       54                       61
##             3TC AZT NVP            3TC AZT EFV
##                      284                      421
```

From this, we can see that the only regimens that were prescribed more than 100 times were "3TC AZT NVP" and "3TC AZT EFV".

Turning this into a data frame is as simple as a call to data.frame, using all_drugs as a set of column labels:

1.9. The dataset haart2.csv contains a few additional observations for the same study. Import these and append them to your master dataset (if you were smart about how you coded the previous steps, cleaning the additional observations should be easy!). Show the first five records and the last five records of the complete (and clean) data set.

```r
haart <- data.frame(read.csv("https://raw.githubusercontent.com/fonnesbeck/Bios6301/master/datasets/haa
haart2 <- data.frame(read.csv("https://raw.githubusercontent.com/fonnesbeck/Bios6301/master/datasets/ha
haart.merged <- merge(haart, haart2, all = TRUE)

# convert date columns
haart.merged[, "init.date"] <- as.POSIXct(haart.merged[, "init.date"], format = "%m/%d/%y")
haart.merged[, "last.visit"] <- as.POSIXct(haart.merged[, "last.visit"], format = "%m/%d/%y")
haart.merged[, "date.death"] <- as.POSIXct(haart.merged[, "date.death"], format = "%m/%d/%y")

# create indicator variable to represent death within 1 year of the initial
# visit
haart.merged$death.within.year <- (difftime(haart.merged$date.death, haart.merged$last.visit,
    units = "days") <= 365)
haart.merged$death.within.year[is.na(haart.merged$death.within.year)] <- 0

# create followup time field
haart.merged <- calculate.followup(haart.merged)

# censor followup time
haart.merged$followup.time <- censor.followup(haart.merged$followup.time)

# add indicator for loss-to-followup
```

```r
haart.merged <- calculate.losstofollowup(haart.merged)

# add indicators for regimen
haart.merged <- create.regimens(haart.merged)

regimen <- matrix(nrow = nrow(haart.merged), ncol = 1)
for (j in 17:34) {
    for (i in 1:nrow(haart.merged)) {
        if (haart.merged[i, j] == TRUE) {
            if (j == 17) {
                regimen[i] <- colnames(haart.merged)[j]
            } else {
                regimen[i] <- paste(regimen[i], colnames(haart.merged)[j])
            }
        }
    }
}
haart.merged <- cbind(haart.merged, regimen)

haart.merged[1:5, ]  #first 5
```

```
##   male age aids cd4baseline    logvl weight hemoglobin    init.reg
## 1    0  18    0          89 5.184231     NA         NA 3TC,AZT,EFV
## 2    0  18    0         280       NA 52.164         11 3TC,AZT,EFV
## 3    0  18    0         431 5.342423 58.000         NA 3TC,AZT,NVP
## 4    0  19    0          51 5.618615 48.600         NA 3TC,AZT,NVP
## 5    0  19    0         180 4.121330     NA         NA 3TC,AZT,NVP
##    init.date last.visit death date.death death.within.year followup.time
## 1 2003-11-03 2006-04-12     0       <NA>                 0           365
## 2 2004-02-19 2008-03-14     0       <NA>                 0           365
## 3 2007-03-13 2007-03-13     0       <NA>                 0             0
## 4 2005-12-07 2007-04-17     0       <NA>                 0           365
## 5 2006-09-08 2006-10-15     0       <NA>                 0            37
##   loss.to.followup init.reg_list  3TC  AZT   EFV   NVP   NFV   ABC   LPV
## 1                0 3TC, AZT, EFV TRUE TRUE  TRUE FALSE FALSE FALSE FALSE
## 2                0 3TC, AZT, EFV TRUE TRUE  TRUE FALSE FALSE FALSE FALSE
## 3                1 3TC, AZT, NVP TRUE TRUE FALSE  TRUE FALSE FALSE FALSE
## 4                0 3TC, AZT, NVP TRUE TRUE FALSE  TRUE FALSE FALSE FALSE
## 5                1 3TC, AZT, NVP TRUE TRUE FALSE  TRUE FALSE FALSE FALSE
##     RTV   D4T   DDI   IDV   SQV   T20   FPV   TDF   ATV   FTC   DDC
## 1 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 2 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 3 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 4 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 5 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
##       regimen
## 1 3TC AZT EFV
## 2 3TC AZT EFV
## 3 3TC AZT NVP
## 4 3TC AZT NVP
## 5 3TC AZT NVP
```

```
haart.merged[1000:1004, ]  #last 5
```

```
##      male age aids cd4baseline  logvl  weight hemoglobin      init.reg
## 1000    1  66    0         298 4.09496     NA         NA    3TC,AZT,EFV
## 1001    1  67    0          95      NA 66.6792         16    3TC,AZT,EFV
## 1002    1  69    0          NA      NA      NA         NA 3TC,AZT,RTV,SQV
## 1003    1  80    0         267      NA 53.0712         NA    3TC,AZT,NVP
## 1004    1  89    0           9      NA 43.5456         10    3TC,ABC,AZT
##        init.date last.visit death date.death death.within.year
## 1000 2006-06-08 2007-02-12     0       <NA>                 0
## 1001 2004-02-13 2008-02-21     0       <NA>                 0
## 1002 2006-04-01 2007-09-13     0       <NA>                 0
## 1003 2004-11-08 2006-11-20     1 2006-11-26                 1
## 1004 2004-12-15 2006-04-11     0       <NA>                 0
##      followup.time loss.to.followup      init.reg_list  3TC   AZT   EFV
## 1000      249.0417                 1      3TC, AZT, EFV TRUE  TRUE  TRUE
## 1001      365.0000                 0      3TC, AZT, EFV TRUE  TRUE  TRUE
## 1002      365.0000                 0 3TC, AZT, RTV, SQV TRUE  TRUE FALSE
## 1003      365.0000                 0      3TC, AZT, NVP TRUE  TRUE FALSE
## 1004      365.0000                 0      3TC, ABC, AZT TRUE  TRUE FALSE
##        NVP   NFV   ABC   LPV   RTV   D4T   DDI   IDV   SQV   T20   FPV
## 1000 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 1001 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 1002 FALSE FALSE FALSE FALSE  TRUE FALSE FALSE FALSE  TRUE FALSE FALSE
## 1003  TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 1004 FALSE FALSE  TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
##        TDF   ATV   FTC   DDC     regimen
## 1000 FALSE FALSE FALSE FALSE 3TC AZT EFV
## 1001 FALSE FALSE FALSE FALSE 3TC AZT EFV
## 1002 FALSE FALSE FALSE FALSE 3TC AZT RTV SQV
## 1003 FALSE FALSE FALSE FALSE 3TC AZT NVP
## 1004 FALSE FALSE FALSE FALSE 3TC AZT ABC
```

Question 2

14 points

Use the following code to generate data for patients with repeated measures of A1C (a test for levels of blood glucose).

```
genData <- function(n) {
    if(exists(".Random.seed", envir = .GlobalEnv)) {
        save.seed <- get(".Random.seed", envir= .GlobalEnv)
        on.exit(assign(".Random.seed", save.seed, envir = .GlobalEnv))
    } else {
        on.exit(rm(".Random.seed", envir = .GlobalEnv))
    }
    set.seed(n)
    subj <- ceiling(n / 10)
    id <- sample(subj, n, replace=TRUE)
    times <- as.integer(difftime(as.POSIXct("2005-01-01"), as.POSIXct("2000-01-01"), units='secs'))
    dt <- as.POSIXct(sample(times, n), origin='2000-01-01')
    mu <- runif(subj, 4, 10)
```

```
    a1c <- unsplit(mapply(rnorm, tabulate(id), mu, SIMPLIFY=FALSE), id)
    data.frame(id, dt, a1c)
}
x <- genData(500)
```

Perform the following manipulations: (2 points each)

2.1. Order the data set by id and dt.

```
x <- x[order(x$id,x$dt),]
```

2.2. For each id, determine if there is more than a one year gap in between observations. Add a new row at the one year mark, with the a1c value set to missing. A two year gap would require two new rows, and so forth.

```
#take subset of observations with a parameter for the id
subset.id <- function(id.num)
{
  return(subset(x,x$id==id.num))
}

#find gaps between observation i and the observation below it that are greater than 365 days
#this is where we will need to insert rows
find.gaps <- function(id)
{
  id.vector <- subset.id(id)
  gaps <- matrix(nrow=nrow(id.vector),ncol=1)
  for (i in 1:nrow(id.vector)-1)
  {
    gaps[i] <- as.integer(difftime(id.vector[i+1,2],id.vector[i,2],units="days"))
  }
  return(gaps)
}

#apply the find.gaps function to all ids 1:50
gaps <- NULL
for (id in 1:50)
{
  gaps <- rbind(gaps,find.gaps(id))
}
x <- cbind(x,gaps)

#find rows where the gap (after the observation) > 365
(gaps.positions <- which(abs(x$gaps)>365))
```

```
##  [1]   36  47  49  55  56  69  71  81  87 110 117 119 126 128 135 136 143
## [18] 151 152 159 190 193 196 209 210 215 216 258 267 279 286 294 310 319
## [35] 326 327 336 341 356 360 370 380 381 388 416 420 424 433 436 444 454
## [52] 471 474 481
```

```
#insert rows for gaps of 1 year
for (i in 1:length(gaps.positions))
```

```
{
  row.below <- gaps.positions[i]
  save.above <- x[1:row.below,] #save all of the rows above
  save.below <- x[(row.below+1):nrow(x),] #save all of the rows below
  x[row.below+1,1] <- x[row.below,1]
  x[row.below+1,2] <- x[row.below,2] + days(365)
  x[row.below+1,3] <- NA
  x[row.below+1,4] <- NA
  x <- rbind(save.above,x[row.below+1,],save.below)
  gaps.positions <- gaps.positions + 1 #increment gap indeces to account for new inserted row
}

#apply the function again to allow for two year gaps
gaps2 <- NULL
for (id in 1:50)
{
  gaps2 <- rbind(gaps2,find.gaps(id))
}
x <- cbind(x,gaps2)

(gaps2.positions <- which(abs(x$gaps2)>365)) #find gaps that were greater than 1 year
```

```
## [1] 169 179
```

```
for (i in 1:length(gaps2.positions))
{
  row.below <- gaps2.positions[i]
  save.above <- x[1:row.below,] #save all of the rows above
  save.below <- x[(row.below+1):nrow(x),] #save all of the rows below
  x[row.below+1,1] <- x[row.below,1]
  x[row.below+1,2] <- x[row.below,2] + days(365)
  x[row.below+1,3] <- NA
  x[row.below+1,4] <- NA
  x <- rbind(save.above,x[row.below+1,],save.below)
  gaps2.positions <- gaps2.positions + 1 #increment gap indeces to account for new inserted row
}
```

2.3. Create a new column visit. For each id, add the visit number. This should be 1 to n where n is the number of observations for an individual. This should include the observations created with missing a1c values.

```
#function that will create a visit field for the parameterized id
count.visits <- function(id)
{
  id.vector <- subset.id(id)
  visit.no <- matrix(nrow=nrow(id.vector),ncol=1)
  for (i in 1:nrow(visit.no))
  {
    visit.no[i] <- i
  }
  return(visit.no)
}
```

```
#run the count.visits function on all ids to create a complete column for visits
visit <- NULL
for (id in 1:50)
{
  visit <- rbind(visit,count.visits(id))
}

#add the visit column to the original x dataframe
x <- cbind(x,visit)
```

2.4. For each id, replace missing values with the mean a1c value for that individual.

```
#function that will find NA in the a1c field and replace them with that id's mean a1c
replace.na <- function(id)
{
  id.vector <- subset.id(id)
  a1c.mean <- mean(id.vector$a1c, na.rm=TRUE)
  new.a1c <- matrix(nrow=nrow(id.vector),ncol=1)
  for (i in 1:nrow(id.vector))
  {
    if (is.na(id.vector[i,3])==TRUE)
    {
      new.a1c[i] <- a1c.mean
    }
    else
    {
      new.a1c[i] <- id.vector[i,3]
    }
  }
  return(new.a1c)
}

#run the replace.na function for every id from 1:50
a1c.replaced <- NULL
for (id in 1:50)
{
  a1c.replaced <- rbind(a1c.replaced,replace.na(id))
}


x <- cbind(x,a1c.replaced)
```

2.5. Print mean a1c for each id.

```
id.means <- function(id)
{
  return(mean(subset.id(id)$a1c.replaced))
}

id <- seq(1:50)

indiv.means <- lapply(id,id.means)
```

```
(avg.a1c.by.id <- cbind(id,indiv.means))
```

```
##          id indiv.means
##  [1,]  1  4.063372
##  [2,]  2  7.544643
##  [3,]  3  6.75764
##  [4,]  4  3.892127
##  [5,]  5  9.512311
##  [6,]  6  7.555965
##  [7,]  7  9.161686
##  [8,]  8  7.189064
##  [9,]  9  9.283873
## [10,] 10 7.975217
## [11,] 11 6.917562
## [12,] 12 7.034021
## [13,] 13 9.145282
## [14,] 14 6.623756
## [15,] 15 8.012406
## [16,] 16 4.222158
## [17,] 17 3.996034
## [18,] 18 9.164873
## [19,] 19 5.50721
## [20,] 20 3.726675
## [21,] 21 8.140939
## [22,] 22 5.637501
## [23,] 23 7.366889
## [24,] 24 7.439316
## [25,] 25 6.877135
## [26,] 26 6.556759
## [27,] 27 4.926457
## [28,] 28 7.433917
## [29,] 29 4.508086
## [30,] 30 6.045577
## [31,] 31 7.116586
## [32,] 32 6.568791
## [33,] 33 6.494069
## [34,] 34 6.768615
## [35,] 35 8.4767
## [36,] 36 9.60441
## [37,] 37 9.606253
## [38,] 38 5.355979
## [39,] 39 6.917013
## [40,] 40 9.530136
## [41,] 41 9.802424
## [42,] 42 3.89177
## [43,] 43 6.095849
## [44,] 44 9.09167
## [45,] 45 6.737204
## [46,] 46 9.621763
## [47,] 47 9.231489
## [48,] 48 6.4046
## [49,] 49 6.096076
## [50,] 50 8.962319
```

2.6. Print total number of visits for each id.

```
total.visits <- function(id)
{
  return(nrow(subset.id(id)))
}
id <- seq(1:50)
total.visits <- lapply(id,total.visits)
(total.visits.by.id <- cbind(id,total.visits))
```

```
##       id total.visits
## [1,] 1  11
## [2,] 2  20
## [3,] 3  14
## [4,] 4  12
## [5,] 5  14
## [6,] 6  10
## [7,] 7  9
## [8,] 8  12
## [9,] 9  11
## [10,] 10 12
## [11,] 11 10
## [12,] 12 10
## [13,] 13 8
## [14,] 14 12
## [15,] 15 8
## [16,] 16 9
## [17,] 17 12
## [18,] 18 10
## [19,] 19 10
## [20,] 20 9
## [21,] 21 10
## [22,] 22 8
## [23,] 23 8
## [24,] 24 15
## [25,] 25 12
## [26,] 26 14
## [27,] 27 11
## [28,] 28 14
## [29,] 29 10
## [30,] 30 7
## [31,] 31 11
## [32,] 32 5
## [33,] 33 8
## [34,] 34 12
## [35,] 35 11
## [36,] 36 9
## [37,] 37 17
## [38,] 38 15
## [39,] 39 8
## [40,] 40 7
## [41,] 41 17
## [42,] 42 14
## [43,] 43 11
```

```
## [44,] 44 11
## [45,] 45 14
## [46,] 46 9
## [47,] 47 12
## [48,] 48 11
## [49,] 49 12
## [50,] 50 10
```

2.7. Print the observations for id = 15.

```
subset.id(15)
```

```
##          id                dt      a1c gaps gaps2 visit a1c.replaced
## 11       15 2000-04-30 00:34:50 7.527105  262   262     1     7.527105
## 406      15 2001-01-17 21:11:02 5.898371   97    97     2     5.898371
## 306      15 2001-04-25 06:23:05 8.566593  772   365     3     8.566593
## 484      15 2002-04-25 06:23:05       NA   NA   407     4     8.012406
## 4841     15 2003-04-25 06:23:05       NA   NA   365     5     8.012406
## 48411    15 2003-06-06 14:06:00 9.133769  441   365     6     9.133769
## 263      15 2004-06-05 14:06:00       NA   NA    76     7     8.012406
## 2631     15 2004-08-20 17:47:11 8.936190   NA    NA     8     8.936190
```

Question 3

10 points

Import the addr.txt file from the GitHub repository. This file contains a listing of names and addresses (thanks google). Parse each line to create a data.frame with the following columns: lastname, firstname, streetno, streetname, city, state, zip. Keep middle initials or abbreviated names in the firstname column. Print out the entire data.frame.

```r
addr <- read.delim("https://raw.githubusercontent.com/fonnesbeck/Bios6301/master/datasets/addr.txt",
    stringsAsFactors = FALSE, head = FALSE)

find.spaces <- function(textRow) {
    spaces <- c(unlist(gregexpr(" {2,}", textRow)), nchar(textRow))
    return(spaces)  #return a vector of values where spaces of length 2+ are located in the textRow
}

trim <- function(x) gsub("^\\s+|\\s+$", "", x)
# credit:
# http://stackoverflow.com/questions/2261079/how-to-trim-leading-and-trailing-whitespace-in-r

findNumbers <- function(textRow) {
    numberPos <- NULL
    for (i in 1:nchar(textRow)) {
        if ((substr(textRow, i, i) %in% seq(0, 9)) == TRUE) {
            numberPos <- c(numberPos, i)
        }
    }
    return(numberPos)
}
```

15

```r
lastname <- vector()
firstname <- vector()
streetno <- vector()
streetname <- vector()
city <- vector()
state <- vector()
zip <- vector()

fixText <- function(text) {
    for (i in 1:nrow(text)) {
        row.spaces <- find.spaces(text[i, ])
        lastname[i] <- substr(text[i, ], 1, (row.spaces[1] - 1))
        firstname[i] <- trim(substr(text[i, ], (row.spaces[1] + 1), (row.spaces[2] -
            1)))
        streetno.pos <- findNumbers(substr(text[i, ], (row.spaces[2] + 1), (row.spaces[3] -
            1)))  #find indices of the numbers in the street address
        # some addresses contain numerical street names, so streetno.pos can be
        # longer than just the first 3-4 streetno values
        if (length(streetno.pos) > 4) {
            streetno[i] <- substr(substr(text[i, ], (row.spaces[2] + 1), (row.spaces[3] -
                1)), min(streetno.pos), (min(streetno.pos) + 3))
            streetname[i] <- trim(substr(substr(text[i, ], (row.spaces[2] +
                1), (row.spaces[3] - 1)), (min(streetno.pos) + 4), row.spaces[3] -
                1))
        } else {
            streetno[i] <- substr(substr(text[i, ], (row.spaces[2] + 1), (row.spaces[3] -
                1)), min(streetno.pos), max(streetno.pos))
            streetname[i] <- trim(substr(substr(text[i, ], (row.spaces[2] +
                1), (row.spaces[3] - 1)), (max(streetno.pos) + 1), row.spaces[3] -
                1))
        }

        city[i] <- trim(substr(text[i, ], row.spaces[3], row.spaces[4]))
        state[i] <- trim(substr(text[i, ], row.spaces[4], row.spaces[5]))
        zip[i] <- trim(substr(text[i, ], row.spaces[5], row.spaces[6]))
    }
    zip <- sub("O", "0", zip)  #replace mistaken Os in zipcodes with 0s
    return(cbind(lastname, firstname, streetno, streetname, city, state, zip))
}

(addr <- data.frame(fixText(addr)))
```

```
##         lastname  firstname streetno           streetname        city state
## 1          Bania  Thomas M.      725    Commonwealth Ave.      Boston    MA
## 2        Barnaby      David      373       W. Geneva St.    Wms. Bay    WI
## 3         Bausch       Judy      373       W. Geneva St.    Wms. Bay    WI
## 4         Bolatto    Alberto      725    Commonwealth Ave.      Boston    MA
## 5       Carlstrom       John      933        E. 56th St.     Chicago    IL
## 6      Chamberlin Richard A.      111         Nowelo St.        Hilo    HI
## 7          Chuss       Dave     2145        Sheridan Rd    Evanston    IL
## 8          Davis      E. J.      933        E. 56th St.     Chicago    IL
## 9          Depoy     Darren      174       W. 18th Ave.    Columbus    OH
## 10        Griffin       Greg     5000        Forbes Ave. Pittsburgh    PA
```

```
## 11    Halvorsen       Nils  933        E. 56th St.    Chicago  IL
## 12       Harper         Al  373      W. Geneva St.   Wms. Bay  WI
## 13        Huang     Maohai  725 W. Commonwealth Ave.    Boston  MA
## 14      Ingalls   James G.  725 W. Commonwealth Ave.    Boston  MA
## 15      Jackson   James M.  725 W. Commonwealth Ave.    Boston  MA
## 16      Knudsen      Scott  373      W. Geneva St.   Wms. Bay  WI
## 17        Kovac       John 5640       S. Ellis Ave.    Chicago  IL
## 18    Landsberg      Randy 5640       S. Ellis Ave.    Chicago  IL
## 19           Lo  Kwok-Yung 1002       W. Green St.     Urbana  IL
## 20  Loewenstein Robert F.  373      W. Geneva St.   Wms. Bay  WI
## 21        Lynch       John 4201        Wilson Blvd  Arlington  VA
## 22      Martini       Paul  174       W. 18th Ave.   Columbus  OH
## 23        Meyer    Stephan  933        E. 56th St.    Chicago  IL
## 24       Mrozek       Fred  373      W. Geneva St.   Wms. Bay  WI
## 25      Newcomb       Matt 5000        Forbes Ave. Pittsburgh  PA
## 26        Novak      Giles 2145        Sheridan Rd   Evanston  IL
## 27       Odalen      Nancy  373      W. Geneva St.   Wms. Bay  WI
## 28       Pernic       Dave  373      W. Geneva St.   Wms. Bay  WI
## 29       Pernic        Bob  373      W. Geneva St.   Wms. Bay  WI
## 30     Peterson    Jeffrey 5000        Forbes Ave. Pittsburgh  PA
## 31        Pryke       Clem  933        E. 56th St.    Chicago  IL
## 32       Rebull      Luisa 5640       S. Ellis Ave.    Chicago  IL
## 33    Renbarger     Thomas 2145        Sheridan Rd   Evanston  IL
## 34      Rottman        Joe 8730 W. Mountain View Ln  Littleton  CO
## 35    Schartman      Ethan  933        E. 56th St.    Chicago  IL
## 36        Spotz        Bob  373      W. Geneva St.   Wms. Bay  WI
## 37        Thoma       Mark  373      W. Geneva St.   Wms. Bay  WI
## 38       Walker      Chris  933      N. Cherry St.     Tucson  AZ
## 39      Wehrer     Cheryl 5000        Forbes Ave. Pittsburgh  PA
## 40        Wirth      Jesse  373      W. Geneva St.   Wms. Bay  WI
## 41       Wright       Greg  791 Holmdel-Keyport Rd.    Holmdel  NY
## 42      Zingale    Michael 5640       S. Ellis Ave.    Chicago  IL
##          zip
## 1      02215
## 2      53191
## 3      53191
## 4      02215
## 5      60637
## 6      96720
## 7  60208-3112
## 8      60637
## 9      43210
## 10     15213
## 11     60637
## 12     53191
## 13     02215
## 14     02215
## 15     02215
## 16     53191
## 17     60637
## 18     60637
## 19     61801
## 20     53191
## 21     22230
```

```
## 22        43210
## 23        60637
## 24        53191
## 25        15213
## 26 60208-3112
## 27        53191
## 28        53191
## 29        53191
## 30        15213
## 31        60637
## 32        60637
## 33 60208-3112
## 34        80125
## 35        60637
## 36        53191
## 37        53191
## 38        85721
## 39        15213
## 40        53191
## 41 07733-1988
## 42        60637
```

Question 4

2 points

The first argument to most functions that fit linear models are formulas. The following example defines the response variable death and allows the model to incorporate all other variables as terms. . is used to mean all columns not otherwise in the formula.

```
url <- "https://github.com/fonnesbeck/Bios6301/raw/master/datasets/haart.csv"
haart_df <- read.csv(url)[,c('death','weight','hemoglobin','cd4baseline')]
coef(summary(glm(death ~ ., data=haart_df, family=binomial(logit))))
```

```
##                  Estimate  Std. Error   z value      Pr(>|z|)
## (Intercept)   3.576411744 1.226870535  2.915069 0.0035561039
## weight       -0.046210552 0.022556001 -2.048703 0.0404911395
## hemoglobin   -0.350642786 0.105064078 -3.337418 0.0008456055
## cd4baseline   0.002092582 0.001811959  1.154872 0.2481427160
```

Now imagine running the above several times, but with a different response and data set each time. Here's a function:

```
myfun <- function(dat, response) {
  form <- as.formula(response ~ .)
  coef(summary(glm(form, data=dat, family=binomial(logit))))
}
```

Unfortunately, it doesn't work. tryCatch is "catching" the error so that this file can be knit to PDF.

```
tryCatch(myfun(haart_df, death), error = function(e) e)
```

```
## <simpleError in eval(expr, envir, enclos): object 'death' not found>
```

18

What do you think is going on? Consider using debug to trace the problem.

The tryCatch error message reads: <simpleError in eval(expr, envir, enclos): object 'death' not found>, so I wondered if the problem could be in the way that the "response" variable was included in the parameters. To test my theory, I ran the line again with a minor correction:

```
tryCatch(myfun(haart_df, haart_df$death), error = function(e) e)
```

```
## Warning: glm.fit: algorithm did not converge
```

```
##                  Estimate   Std. Error        z value  Pr(>|z|)
## (Intercept) -2.656607e+01 115935.1724 -2.291459e-04 0.9998172
## death        5.313214e+01  69028.4183  7.697140e-04 0.9993859
## weight      -4.499694e-15   1939.0571 -2.320558e-18 1.0000000
## hemoglobin   5.124642e-14   9774.8190  5.242697e-18 1.0000000
## cd4baseline  1.830771e-16    184.0846  9.945271e-19 1.0000000
```

Hooray! Now the function runs completely. Although, the algorithm does not converge, but that is more an artifact of the data than the code written above.

Bonus

5 bonus points

Create a working function.

```r
myfun <- function(dat, response) {
  attach(dat, warn.conflicts = FALSE)
  response.name <- deparse(substitute(response))
  df.name <- deparse(substitute(dat))
  reponse.df <- paste(df.name,response.name, sep="$")
  formula <- paste(reponse.df," ~ .", sep="")
  #by attaching the dataframe names we can call response without an identifier
  #form <- as.formula(response ~ .)
  print(coef(summary(glm(formula, data=dat, family=binomial(logit)))))
  detach(dat)
}

tryCatch(myfun(haart_df, death), error = function(e) e)
```

```
##                  Estimate  Std. Error   z value     Pr(>|z|)
## (Intercept)  3.576411744 1.226870535  2.915069 0.0035561039
## weight      -0.046210552 0.022556001 -2.048703 0.0404911395
## hemoglobin  -0.350642786 0.105064078 -3.337418 0.0008456055
## cd4baseline  0.002092582 0.001811959  1.154872 0.2481427160
```